

Compiler-Enabled Cache Management for Pointer-Intensive Codes

Yao Guo, Saurabh Chheda, Csaba Andras Moritz

Department of Electrical & Computer Engineering
University of Massachusetts, Amherst

1

Motivation

- Compiler-managed caches have been shown to be successful in energy optimizations
 - Real programs often contain pointer intensive parts
 - Precise pointer related analysis is difficult or often not possible for large/complex programs
- Question: can we do better with pointers than just ignoring them?
- Our proposal
 - Speculative Pointer and Distance Analysis
 - Architectural support

2

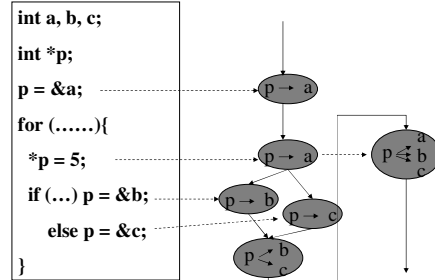
Background on Pointer Analysis

- **Pointer Analysis** determines the set of possible values (called location sets) for each pointer variable in a program
 - Type based (e.g. Steensgaard '96)
 - global results (less precise)
 - Dataflow based (e.g. Emami '94, Wilson '95)
 - a result for each program point
 - using points-to graphs
 - most precise
- Our Approach is based on Points-to Graphs

3

Points-to Graphs

- Dataflow based pointer analysis



4

Efficient Management of Pointer Accesses

- What information do we need to know about pointers in compiler managed caches?
- How do we perform pointer analysis to gather this information?
- Challenges
 - Complex/large programs, precision, static libraries, undefined pointers, recursion

5

Our Approach

- Speculative Pointer and Distance Analysis
 - Speculate when you don't have provable information
 - Speculate to reduce complexity or to complete analysis
 - Speculate on pointer access patterns to evaluate strides (locality)
 - Speculate to improve precision of analysis
 - precision is the width of location set
- Context is compiler managed caches
 - Architectural support to validate
 - Compile-time tradeoffs: degree of speculation

6

Pointer Information Extracted

- *Speculative points-to information*
 - Context and flow sensitive location sets for pointers
- *Frequencies* of locations within location sets
 - We can remove less frequent locations
- Determine *if critical* (e.g., loop-based) or not
 - Focus on more critical pointer variables
- Determine *locality*
 - Pointers change during each loop iteration
 - Stride can be determined with distance analysis
 - Various degrees of confidence can be established
- Determine when a pointer access *is irregular*
 - Negative result is actually useful

7

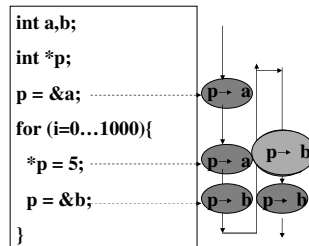
Speculate to Complete Analysis

- Applications have static libraries and sources may not be always available
 - Ignore or speculate on their impact on location sets
- Undefined pointers
 - ignore or treat as no information available (e.g., map to dynamic cache management)
- Recursive procedure calls
 - Speculatively stops after a fixed # of iterations

8

Speculative Optimization for Precision

- Optimize pointer location set widths in loop bodies speculatively
- Remove unlikely or infrequent location sets
 - Improves ability to manage access statically in our context
 - Improves precision for the typical case



9

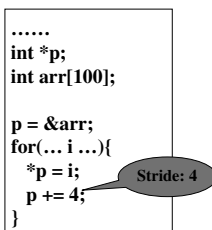
Speculative Distance Analysis

- Objective
 - Capture strides in pointer accesses within loops
 - Determine if there is locality (e.g., if strides are small)
 - Fixed strides can be easily determined
 - Program variable strides can be determined speculatively

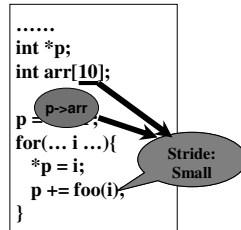
10

Speculative Distance Analysis Examples

Fixed Stride



Variable Stride



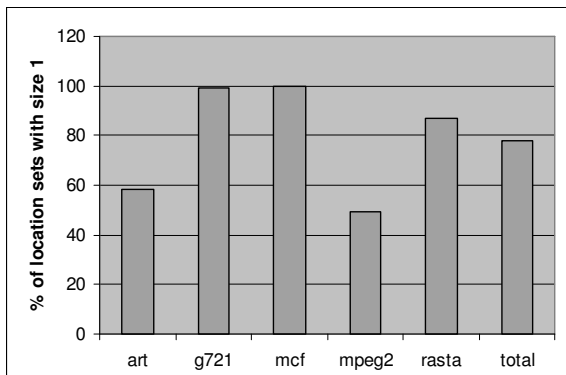
11

Applicability in Cache Management

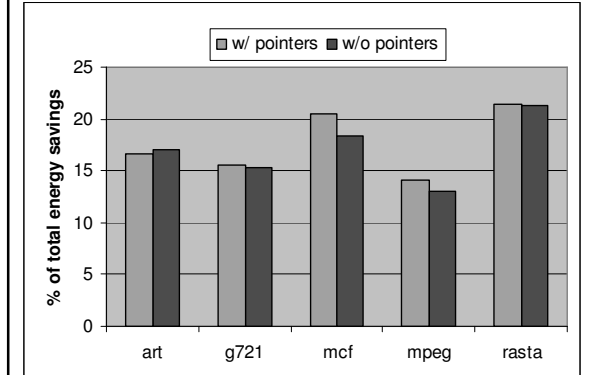
- Cool-Mem (our previous work, ASPLOS 2002)
- Combines static and dynamic memory management
 - fully static (most energy efficient)
 - statically speculative (almost as energy efficient)
 - conventional (least energy efficient)
- Improves energy efficiency by reducing redundancy
- More statically managed accesses reduce energy consumption

12

Results – Location Sets Distribution



Results – Energy Savings Comparison



Conclusions

- **Traditional pointer analysis**
 - Based on compile time provable information
 - Stops analysis if that cannot be guaranteed
 - Used in many performance optimization techniques
- **Speculative pointer and distance analysis**
 - Always completes analysis without restrictions
 - Extracts precise information for our purposes
 - Targets compiler-enabled memory systems
- **Preliminary results**
 - Saves energy for pointer-intensive programs

15